

Digital KAM Blueprint

A living key-account platform for Psiog. Account plans move out of static Excel workbooks into a versioned, scored and AI-assisted system. Account teams, BD and the Executive team each get their own view, and every account gets a recommended next best action.


- Azure-hosted · 7 Azure services

4 AI providers · 6 agents
- All 10 acceptance criteria mapped

Synthetic data only

01 · ARCHITECTURE

How the pieces connect



The browser only talks to the KAM API. The API owns the data and calls the AI Orchestrator, so LLM keys never reach the client. The orchestrator picks a model per task and falls back to another provider when one hits a rate limit. The API writes uploaded workbooks to Blob Storage, and that write triggers the import Function. Dashed boxes and lines are external inputs and secret reads. Every Azure service reads its secrets from Key Vault through managed identity. The same REST API is published for BD and Executive consumers.

Three paths the demo should show

AC 5 • IMPORT

Excel workbook to living plan

1. User uploads a KAM .xlsx. The API stores it in Blob Storage.
2. A blob-trigger Function parses it with openpyxl.
3. Gemini maps non-standard cells to the KAM schema and gives a confidence score per field.
4. A deterministic validator checks required fields, RACI, A/B/C importance and dates.
5. The validation report is shown in the UI. The user confirms and the plan lands in Azure SQL as version 1.

AC 8 • AI

Next best action

1. A nightly Function, or the “Suggest” button, calls the orchestrator.
2. Claude calls tools: plan gaps, scores, stakeholder map, overdue actions.
3. It returns JSON: what to do, whose help, which channel, why, at-risk flag, source sections.
4. The output is schema-checked and logged with model, prompt and tokens.
5. The user accepts it and it becomes a tracked action linked to its section.

AC 2 • 3 • 4 • CADENCE

Continuous planning loop

1. Every save creates a new version through SQL temporal tables.
2. The scoring engine recalculates opportunity, brickwall, checklist and perception scores.
3. Timer Functions flag stale sections and overdue actions, then raise in-app reminders.
4. A nightly extract moves data to Databricks, where score trends are built for the MAC and Power BI views.

What each tool does for Digital KAM

Everything comes from the pSpark approved list. Items marked clear first are frameworks or libraries to run past the core committee before 10:30.

TOOL	ROLE	WHAT IT DOES FOR OUR USE CASE
AI ASSISTANTS AND MODELS		
Claude API	Runtime reasoning agents	Runs the next best action, risk explanation and executive brief agents. Sonnet 5 handles reasoning and Haiku 4.5 handles bulk classification. Prompt caching and nightly batching keep spend within the \$25 credit.
Claude seats + Claude Code	Build-time copilot	Used for planning, schema design, code, tests and docs. Subagents generate tests and review PRs, and the GitHub connector works on issues. Seats on the second org cover build work so the API credit is kept for runtime.
Gemini API	Document understanding	Reads a whole KAM workbook in one long-context call, maps it to our schema and explains mismatches for the validation report. It is also the fallback model when Claude hits a limit.
Groq API	Low-latency copilot	Responds while users edit sections. It drafts 3-year and 1-year vision, suggests strategies from the menu and rewrites challenges. Speed matters here more than depth.
Ollama (local)	Private, offline model	Generates fictional accounts, stakeholders and opportunities so no customer data is used. It also scores AI outputs against a golden test set (LLM-as-judge) and works offline in dev.
APPLICATION		
React + TypeScript (Vite) clear first	Web portal	Plan wizard with completion status per section. Account, MAC, NBD, EBD and Admin views. The NBA panel sits on each account.
Recharts clear first	In-app charts	Score trend lines, brickwall radar, opportunity matrix (potential × effort) and

TOOL	ROLE	WHAT IT DOES FOR OUR USE CASE
		portfolio health tiles.
Python FastAPI clear first	REST API	Accounts, plans, scores and actions endpoints, with OpenAPI/Swagger docs generated automatically (AC 6). Hosts the scoring engine and RBAC. Python is shared with Functions and the AI code.
openpyxl + pandas clear first	Excel round trip	Parses legacy workbooks. Writes plans back in the same structure and builds the BD and Executive export packs.
AZURE (DEPLOYMENT BASE)		
Static Web Apps	Frontend hosting	Serves the SPA with built-in auth and role routing for Account, Group and Executive access.
App Service	API hosting	Runs the KAM API, the system's single front door for the UI and external consumers.
Functions	Background jobs	Review-cadence reminders, stale and overdue flags, blob-triggered Excel import, nightly NBA batch and Databricks extracts.
Container Apps	AI orchestrator	Runs the agent service as a container built with Podman. It scales on its own and keeps LLM rate limits and retries away from the core API.
Azure SQL	System of record	Relational KAM model. Temporal tables give version history and change diffs for free. Master data tables make offerings, strategies, weights and captains configurable (AC 9).
Blob Storage	Files	Raw uploaded workbooks, validation reports, generated export packs and Databricks extract files.

TOOL	ROLE	WHAT IT DOES FOR OUR USE CASE
Key Vault	Secrets	Claude, Gemini and Groq keys and DB connection strings, read through managed identity. No secrets in the repo.
DATA AND BI		
Databricks (free)	Analytics platform	Bronze, silver and gold tables. Builds score trends across versions, action velocity and at-risk features that feed the risk agent.
Power BI Desktop	Executive dashboard	Portfolio summary, account health, NBD pipeline and EBD cross-sell by offering. The .pbix is submitted with the artifacts.
Metabase clear first	Self-serve (optional)	Read-only ad-hoc questions for BD on Azure SQL. Hosted on Container Apps, and only built if time allows.
DELIVERY, QA AND DESIGN		
GitHub	SDLC backbone	Repo, Issues and a Projects board mapped to the 10 acceptance criteria, PR reviews, and Actions CI/CD deploying to Azure. Also holds the submitted artifacts.
Podman · Azure CLI · GitHub CLI	Local and scripted ops	Local containers for the API and orchestrator. A provisioning script that anyone can re-run for the resource group. Repo automation.
pytest · Vitest clear first	Quality	Unit tests for the scoring engine, API contract tests, UI tests and an AI golden-set eval. Test reports are committed for offline evaluation.
draw.io · Vertabelo · Excalidraw	Design artifacts	Architecture diagram, ERD for the KAM schema and wireframes of the plan wizard.

Six agents, each with one job

Each agent returns structured JSON checked against a schema. A person accepts every AI suggestion before it changes a plan.

Claude Sonnet 5 • tool use

Next Best Action

Recommends what to do next for each account, whose help to seek, which channel to use and why. It works from plan gaps, scores, the stakeholder map and overdue actions.

out → action, helper, channel, rationale, at_risk, source_sections

Gemini • long context

Workbook Interpreter

Maps legacy Excel layouts to the KAM schema, gives a confidence score per field and explains what it could not map.

out → mapped rows + validation findings

Groq • Llama

Plan Copilot

Helps users while they fill in sections. It drafts vision and objectives, suggests strategies and checklist answers, and summarises stakeholder notes.

out → draft text the user edits or rejects

Claude Haiku 4.5 • nightly batch

Risk Sentinel

Reads score trends, stale sections and overdue actions from Databricks gold tables, then flags at-risk accounts with a plain-language reason.

out → risk level, drivers, flagged sections

Claude Sonnet 5

Executive Briefer

Writes the portfolio narrative for the Executive export pack and the MAC view: what moved, where it is at risk and where to invest.

out → briefing sections with cited accounts

Ollama • local model

Synthetic Data & Evaluator

Creates realistic fictional accounts for the demo and tests. Scores the other agents' outputs against a rubric in the test suite.

out → seed data, eval scores per agent

- Synthetic or public data only. No customer data or Psiog IP goes to any model.
- Every AI call is logged with agent, model, tokens and outcome in an AI audit table.
- Keys stay in Key Vault and are only called server-side through the orchestrator.
- A token budget per agent, response caching and a provider fallback keep us running if one model hits its rate limit.

Acceptance criteria to components

AC	CRITERION	WHERE IT IS BUILT
1	Digitised account plan	SPA wizard with 8 guided sections and completion % per section. Section tables in Azure SQL.
2	Scoring	API scoring engine. Criteria and weights read from master data, so changes need no deploy.
3	Continuous planning	Temporal tables for versions and diffs. Timer Functions for cadence, stale and overdue flags. Databricks for score trends.
4	Action tracking	Actions table with owner, due date, status and a link to its source section. Status board in the SPA.
5	Excel import and export	Blob-trigger Function, Gemini mapping, deterministic validator and report. openpyxl export in the same structure.
6	REST API and tailored exports	FastAPI with OpenAPI docs. BD pack (opportunities, stakeholders, actions) and Executive pack (summary, health, risks).
7	Views	Account, MAC, NBD and EBD screens in the SPA. Executive portfolio in Power BI.
8	AI next best action	AI Orchestrator with the Next Best Action and Risk Sentinel agents.
9	Configurable	Admin screens for offerings, strategy menu, scoring criteria, account groups and captains.
10	Access	JWT roles for Account, Group and Executive access, enforced by row scoping in the API.

06 · TEAM

Suggested split for five

- **Architect (you):** architecture, AI orchestrator and agents, reviewing AI output, committee liaison.
- **Backend:** SQL schema, FastAPI, scoring engine, RBAC.
- **Frontend:** SPA views, plan wizard, in-app charts.

07 · BEFORE 10:30

Clear with the core committee

- Frameworks: React + TypeScript (Vite), FastAPI, Recharts, openpyxl and pandas, pytest and Vitest.
- Whether the \$25 Claude credit can be used as an API key for runtime agents.

- **Integration and data:** Excel import and export, Functions, Databricks, Power BI.
- **QA and docs:** test plan, pytest and Vitest, AI golden set, demo script, README for offline evaluation.
- Running Metabase in a container on Container Apps.
- Entra ID sign-in for Static Web Apps, or app-issued JWTs instead.
- Email reminders need a service outside the list. Plan for in-app reminders unless one is approved.

Use case questions close at lunch, so raise open questions about the KAM brief in the morning.